

# Query Auto-Completion System

## End-to-End Pipeline Design

# Stage 1: Feature Engineering

## Input Data:

- df\_pool: 150K unique queries
- df\_query\_features: catalog\_clicks, orders, volume, catalog\_views

## Processing Steps:

- ① Merge pool with query features on 'query' column
- ② Create lowercase version: query\_lower
- ③ Build 3-char prefix index: prefix\_3char
- ④ Normalize features using min-max scaling
- ⑤ Compute popularity score with weighted formula

## Popularity Score Formula:

$$p = 0.4 \cdot \text{orders\_norm} + 0.3 \cdot \text{clicks\_norm} + 0.2 \cdot \text{volume\_norm} + 0.1 \cdot \text{views\_norm}$$

**Output:** df\_pool\_enhanced with popularity scores

# Stage 2: Dual Encoding Pipeline

## Dense Path: Semantic Embeddings

- ① Load all-MiniLM-L6-v2
- ② Batch encode: 64 queries/batch
- ③ Original: 384-dimensional
- ④ Apply PCA: 384D → 128D
- ⑤ Save as .npy (float32)

## Why PCA?

- Reduces unnecessary dimension.
- Bearable accuracy loss

## Sparse Path: Lexical Matching

- ① Initialize BM25 with HashingVectorizer
- ② Set  $n\_features = 2^{14} = 16384$
- ③ Fit on pool queries (IDF computation)
- ④ Batch transform: 10K queries/batch
- ⑤ Apply sparsification: threshold = 0.01
- ⑥ Save as CSR matrix (.npz)

## Stage 3: Index Construction

### FAISS Dense Index:

- Type: IndexFlatIP (Inner Product for cosine similarity)
- Dimension: 128D
- Vectors: 150,000
- Index size: ~75 MB
- Search complexity:  $O(n)$  but highly optimized with SIMD

### Prefix Lookup Table:

- 3-character prefix mapping
- Example: "nik" → [123, 456, 789, ...]
- Reduces search space: 150K → 1-5K queries
- Stored in DataFrame index for  $O(1)$  lookup

### All Components Serialized:

- BM25 model, PCA transformer, FAISS index, Dense embeddings, Sparse matrix, Pool metadata

# Stage 4: Inference & Scoring

## Query Processing Flow:

```
# 1. Prefix filtering (if len >= 3)
mask = df_pool_enhanced['prefix_3char'] == prefix.lower()[:3]
filtered_indices = mask[mask].index.to_numpy()

# 2. Encode query with both paths
dense_vec = model.encode([prefix])[0]
dense_vec = pca.transform([dense_vec])[0]    # 128D
sparse_vec = bm25.transform_query(prefix)    # BM25 IDF

# 3. Normalize and weight (alpha = 0.6)
hdense = normalize(dense_vec) * alpha
hsparse = normalize(sparse_vec) * (1 - alpha)

# 4. FAISS search on filtered/full pool
dense_scores, indices = faiss_index.search(hdense, top_k * 2)

# 5. Sparse similarity with CSR matrix
sparse_scores = pool_sparse_matrix[indices].dot(hsparse.T).flatten()

# 6. Combine with popularity boost (w = 0.15)
final_scores = dense_scores + sparse_scores + w * popularity[indices]

# 7. Top-K selection with argpartition
top_indices = np.argpartition(final_scores, -k)[-k:]
```

# Performance & Optimizations

## Speed Optimizations:

- **Numba JIT**: Score combination in compiled code
- **Batch encoding**: 5x faster than sequential
- **Prefix filtering**: 30x search space reduction
- **FAISS SIMD**: Hardware-accelerated similarity
- **CSR sparse matrix**: Memory-efficient sparse ops

## Configuration:

- $\alpha = 0.6$  (dense/sparse balance)
- $w = 0.15$  (popularity weight)
- $k = 100$  (candidates)

# Approach-2: Overview & Motivation

## Core Philosophy

Optimize for **speed** over semantic complexity using lightweight, deterministic methods

### Key Advantages:

- Extremely fast - no neural models at inference
- Simple deployment - single Python script
- Explicit typo handling via fuzzy matching
- Low memory footprint

### Expected Limitations:

- Semantic gaps - misses paraphrases
- Lower hit rate on unseen queries
- No learning from query relationships

*Trade-off: Sacrifice semantic understanding for millisecond-level response times*

# Approach-2: Four Core Components

## ① Historical Frequency Lookup

- Direct prefix → query mapping from training data
- Captures actual user selection behavior
- *Example:* prefix 'sho' → ['shoes': 450, 'shorts': 320, 'shoulder bag': 180]

## ② Prefix Index (Inverted Index)

- Maps each possible prefix to all queries starting with that prefix
- Memory-efficient: limit queries stored per prefix
- Enables instant lexical matching

## ③ Fuzzy Matching

- Handles typos and spelling variations
- Computes similarity between input prefix and popular queries
- Threshold-based candidate selection
- *Example:* "shoez" ≈ "shoes" (similarity > 0.8)

## ④ Popularity-Based Backfill

- Fallback mechanism when < 150 candidates found
- Uses most popular queries from entire pool
- Ensures complete 150-query output

# Scoring, Ranking & Implementation

## Weighted Scoring Function

Final candidates scored using weighted combination:

$$\text{Score}(q|p) = w_1 \cdot S_{\text{hist}}(q, p) + w_2 \cdot S_{\text{prefix}}(q, p) + w_3 \cdot S_{\text{fuzzy}}(q, p) + w_4 \cdot S_{\text{pop}}(q)$$

where:


- $S_{\text{hist}}$  = Historical match score (normalized frequency)
- $S_{\text{prefix}}$  = Direct prefix match score (binary or position-based)
- $S_{\text{fuzzy}}$  = Fuzzy similarity score (e.g., Levenshtein-based)
- $S_{\text{pop}}$  = Normalized popularity score from engagement metrics

# Scoring, Ranking & Implementation

## Implementation Strategy

- **Technology:** Single Python script with standard libraries
- **Data structures:** Dictionaries for frequency maps, prefix indices
- **Fuzzy matching:** RapidFuzz or similar lightweight library
- **No neural models loaded** → instant startup, minimal RAM
- **Output:** Top-150 queries ranked by composite score

# Complete Retrieval Pipeline



# Approach 3: Generative LLM (GPT-2 + LoRA)

**Hypothesis:** Treat auto-complete as a text generation task rather than retrieval.

## Model Architecture & Setup

- **Base Model:** GPT-2 Small (124M parameters).
- **Fine-tuning Strategy:** LoRA (Low-Rank Adaptation).
- **Trainable Parameters:**  $\approx 600K$  (0.5% of total).
- **Input Format:** Prefix: <input> \n
- **Target Output:** Completion: <term1>, <term2>...

**Objective:** Test if the model can learn the query distribution and generate valid completions based on prefix alone.

# LLM Approach: Technical Failures

Experimentation revealed three fundamental blockers for production usage.

## ① Inference Latency (Critical)

- **Observed:** 3-4 seconds per request.
- **Cause:** Autoregressive generation is sequential ( $O(N)$ ). And LLM slower comparatively.

## ② Hallucination & Constraints

- **Issue:** Model generates valid English words that are *not* in our specific query pool.
- **Stat:** lot of generations were "out-of-vocabulary" regarding the database.
- **Correction Cost:** Post-filtering invalid terms increases latency further.

## ③ Stochasticity

- Auto-complete requires deterministic ranking based on popularity.
- LLM outputs vary based on temperature/seed (inconsistent user experience).

# Architecture & Model

- **Model:** OpenAI CLIP ViT-B/32
- **Embedding dim:** 512
- **Normalization:** L2 (cosine similarity)

## Data Pipeline:

- Dataset: Amazon Reviews 2023 (All\_Beauty)
- Filter: `images.notna()`, `title.len() < 10`
- Sample: 5000 products
- Batch size: 32

# Image Encoding

```
# Batch processing with CLIP
image_inputs = torch.stack([
    preprocess(img) for img in batch_images
]).to(device)

with torch.no_grad():
    features = model.encode_image(image_inputs)
    # L2 normalization for cosine similarity
    features = features / features.norm(dim=-1, keepdim=True)
```

## Key Operations:

- Image preprocessing: PIL → CLIP transform
- Encoding: ViT-B/32 encoder
- Output: normalized 512-dim vectors

# FAISS Indexing & Search

```
# FAISS Inner Product index
embeddings_np = embeddings.numpy().astype('float32')
dimension = embeddings_np.shape[1] # 512
index = faiss.IndexFlatIP(dimension)
index.add(embeddings_np)

# Text query encoding
text_tokens = clip.tokenize([query]).to(device)
text_features = model.encode_text(text_tokens)
text_features /= text_features.norm(dim=-1, keepdim=True)

# Similarity search
similarities, indices = index.search(query_np, top_k)
```

**Search:** IndexFlatIP (Inner Product = Cosine Sim)

# Comparative Analysis & Conclusion

## Final Decision

- **Reject Approach 3 (LLM):** Too slow, uncontrollable output.
- **Reject Approach 4 (CLIP):** Solves wrong problem (semantic vs lexical), lacks image data.
- **Recommendation:** Proceed with **Approach 1** (BM25 + SBERT Re-ranking) or **Approach 2** (Trie/Frequency) for production.